

# DSC5103 Statistics

Session 7. Tree-based Methods

# Last time

- Subset Selection
  - Subset Selection and Stepwise Selection revisited
  - Choosing the optimal model using Cross-Validation
- Shrinkage Methods (Regularization)
  - Ridge Regression
  - The Lasso
  - Elastic Net

# Linear Model Selection unified

- Best subset selection

$$\min RSS; \text{ subject to } \|\beta\|_0 \leq A$$

- LASSO

$$\min RSS; \text{ subject to } \|\beta\|_1 \leq A$$

- Ridge Regression

$$\min RSS; \text{ subject to } \|\beta\|_2 \leq A$$

# Plan for today


- The Basics of Decision Trees
  - Regression Trees
  - Building and Pruning Trees
  - Classification Trees
- Trees vs. Linear Models
- Advantages and Disadvantages of Trees

# Tree-based Methods

- The idea: to segment/partition the predictor (X) space into a number of simple regions and predict Y based on the regions
- The set of splitting rules used to segment the predictor space can be summarized in a tree, so these types of approaches are known as **decision-tree methods**
  - Regression trees
  - Classification trees

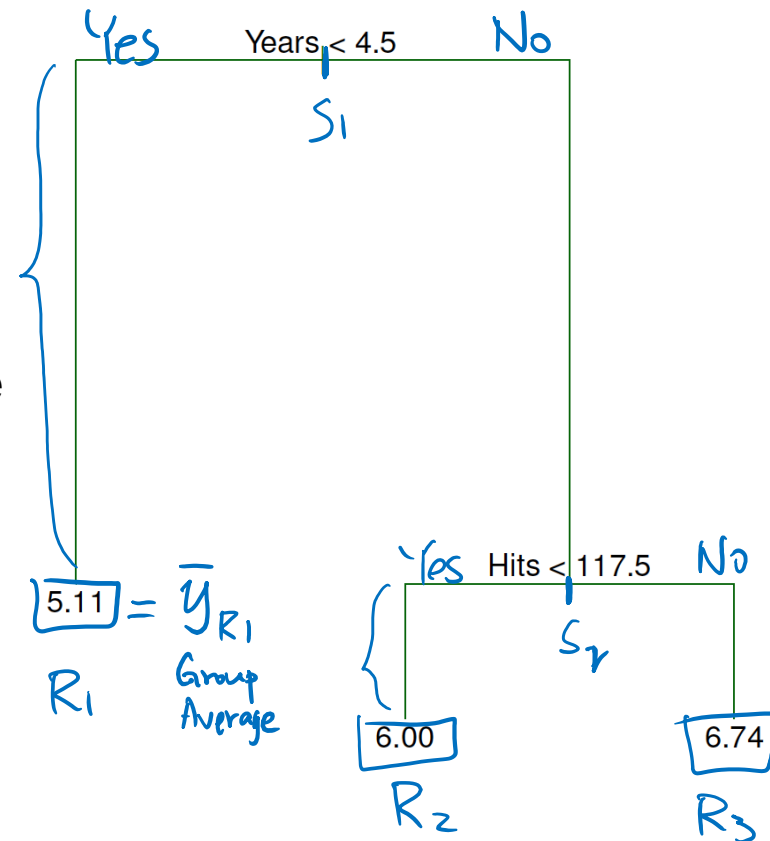
# Example: Hitter's Salary

- Segment on  $X_1$  and  $X_2$  based on  $y$





# Hitter's Salary --- fit the decision tree

- A regression tree:  $\log(\text{salary}) \sim \text{Years} + \text{Hits}$ .
- The tree has two internal nodes and three terminal nodes (leaves).
- At a given internal node, the label indicates the splitting rule.
  - For instance, the split at the top of the tree results in two large branches. The left-hand branch corresponds to  $\text{Years} < 4.5$ , and the right-hand branch corresponds to  $\text{Years} \geq 4.5$ .
- The number in each leaf is the mean of the response for the observations that fall there.



# Hitter's Salary --- visualize the tree

- The tree segments the data into three regions of X space.






# Hitter's Salary --- interpret the tree

- *Years* is the most important factor in determining *Salary*, and players with less experience earn lower salaries than more experienced players.
- Given that a player is less experienced, the number of *Hits* that he made in the previous year seems to play little role in his *Salary*.
- But among players who have been in the major leagues for five or more years, the number of *Hits* made in the previous year does affect *Salary*, and players who made more *Hits* last year tend to have higher salaries.
- Surely an over-simplification, but compared to a regression model, it is easy to display, interpret and explain.

# Hitter's Salary --- predict with the tree

- The tree divides the predictor space into  $J$  distinct and non-overlapping regions,  $R_1, R_2, \dots, R_J$ .
- For every new observation that falls into the region  $R_j$ , we make the same prediction, which is simply the mean of the response values for the training observations in  $R_j$ .



# The Tree Building Process

- The goal: find regions  $R_1, R_2, \dots, R_J$  to minimize RSS:

$$\min_{R_1, R_2, \dots, R_J} \sum_{j=1}^J \sum_{i \in R_j} \underbrace{(y_i - \underbrace{\bar{y}_{R_j}}_{\text{Group Average in } R_j})^2}_{\text{Squared error of } i}$$

$\underbrace{\hspace{10em}}_{\text{total squared error in } R_j}$


$\underbrace{\hspace{15em}}_{\text{total squared error}}$

- In theory, the regions could have any shape. However, we choose to divide the predictor space into high-dimensional rectangles, or boxes, for simplicity and for ease of interpretation of the resulting predictive model.


# The Tree Building Process

- How to find the optimal partition  $R_1, R_2, \dots, R_J$  ?
  - Computationally infeasible to optimize
- CART (Classification And Regression Tree): a top-down greedy approach (“forward stepwise”)
  - it begins at the top of the tree
  - successively splits the one of the regions along one of the  $X$ 's each time
  - each split is indicated via two new branches further down on the tree


# Where to Split?




- Find one predictor  $X_j$  and a cut point  $s$  among all possible values of  $j$  and  $s$  to minimize the RSS.
- The first split is at Years < 4.5
  - Left branch: Years < 4.5
  - Right branch: Years  $\geq 4.5$



# Where to Split?



- Repeat the process, looking for the best predictor and best cut point in order to split the data further so as to minimize the total RSS across all regions
- Instead of splitting the entire predictor space, we split one of the previously identified regions

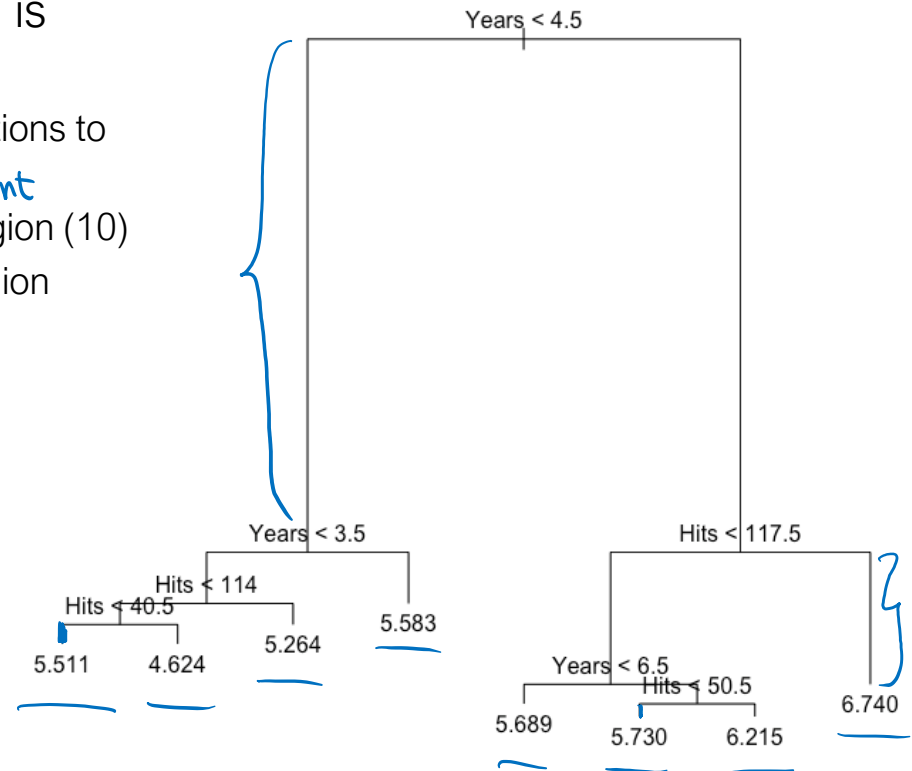


# Where to Split?





- Continues until a stopping criterion is reached
  - mincut: the minimum number of observations to include in either child node (5)
  - minsize: too few data points in a node/region (10)
  - mindev: too small differences within a region
  - ...

- Obtain a tree

- Internal nodes  $\Leftrightarrow$  splits
- Terminal nodes  $\Leftrightarrow$  regions
- Length of branch  $\Leftrightarrow$  reduction in RSS




# Tree Output





# Tree Prediction




# Tree Pruning

- The tree-building process is forward stepwise expansion
  - As the model complexity (**number of leaves**) increases, it tends to overfit the data
- The tree-pruning strategy is backward shrinkage (subset selection for trees)
  - Start with building a large tree  $T_0 \Rightarrow m$  regions  $R_1, \dots, R_m$
  - Prune it backward: for each value of  $\alpha$  (a tuning parameter), find a subtree  $T$  to minimize

$$\sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T| \quad \min \sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \hat{y}_{R_m})^2; \text{ subject to } |T| \leq A$$

- $|T|$  is the number of leaves/regions
- Use cross-validation to find the optimal  $\alpha$  (or tree size  $A$ )



# The Final Tree Algorithm


*enough*

1. Build a large tree using all the training data (`tree()` in R)
2. Run k-fold cross-validation to choose tree size A (or parameter  $\alpha$ ) (`cv.tree()` in R)
  1. In each iteration
    1. use the k-1 training folds to build a large tree
    2. Prune the tree for each tree size A (or  $\alpha$ )
    3. Evaluate MSE on the validation fold for each tree size A (or  $\alpha$ )
  2. Find the best tree size A (or  $\alpha$ )
3. Prune the large tree to the optimal tree size A (`prune.tree()` in R)


\*\* How/why is this different from the subset selection of linear models? \*\*

# Hitter's Salary --- Cross-Validation

- Size=3 seems good




# Hitter's Salary --- Final Pruning



# Classification Trees

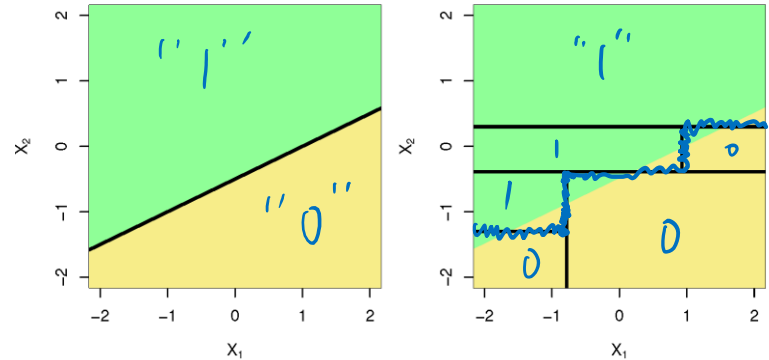
- Very similar to a regression tree, except for
  - Prediction is not based on mean but the proportion of each class in the region (similar to KNN)
  - When splitting a tree, RSS is no longer the right criterion
    - Classification error rate is the natural criterion, but not sensitive enough
      - Also cutoff dependent
    - Deviance
    - Gini index: a measure of total variance across the K classes (“purity”)
$$G = \sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$$
    - Deviance and Gini index are very similar numerically
  - When pruning or optimizing tree size, use deviance / classification error / AUC as the criterion

# Example: Heart data




# Trees vs. Linear Model: Classification Example

- Top row: the true decision boundary is linear
  - Left: linear model (good)
  - Right: decision tree




- Bottom row: the true decision boundary is non-linear
  - Left: linear model
  - Right: decision tree (good)

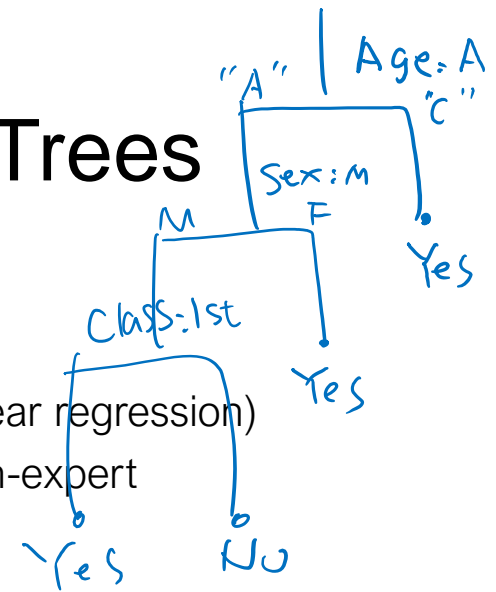





# Trees vs. Linear Model: Regression Example



# Pros and Cons of Decision Trees



- Pros:
  - Very easy to explain to people (probably even easier than linear regression)
  - Can be plotted graphically and easily interpreted even by non-expert
  - Work on both classification and regression problems
  - Capture nonlinear effect (no more transformations on X or Y)
  - Handle categorical predictors (no more dummy variables)
  - Handle interactions (no more \* or : )
  - Handle missing data
- Cons:
  - Trees don't have the same prediction accuracy as some of the more complicated approaches that we examine in this course
  - Final tree is not very stable
  - Computational issue with big categorical variables